

Document Database and MapReduce

Hogeschool Rotterdam
Rotterdam, Netherlands

Lecture topics

- CAP Theorem ACID vs BASE
- Document Databases
- MongoDB
- Map-Reduce
- Summery

Motivation

- As we mentioned before relational database systems are designed to run on a single server
- RDBMS satisfy the ACID rules to provide consistency and availability of the data for the users
- But how do NoSQL databases deal with the data in their implementation?

CAP theorem

- States that it is impossible for a distributed computer system to simultaneously provide all three of the following guarantees:
 - **Consistency** every read receives the most recent write or an error
 - **Availability** every request receives a response, without guarantee that it contains the most recent version of the information
 - **Partition** tolerance (the system continues to operate despite arbitrary partitioning due to network failures)

NoSQL database and CAP theorem

ACID vs BASE

- Basically **A**vailable, **S**oft State and **E**ventual **C**onsistent
- Because of this characteristic the query language must be able to process data saved locally and in a cluster (to be discussed in another slide)

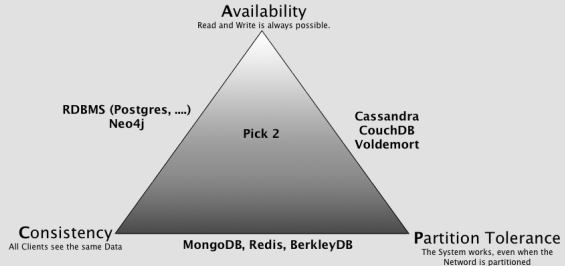


Figure: CAP Theorem

ACID vs BASE

- Relational databases are considered to be structural
- Document databases uses semi-structured formats
- text files such as logs are unstructured

Introduction

- A document database is a nonrelational database that stores data as semi-structured documents such as in XML or JSON formats
- Document databases are free to implement ACID transactions or other characteristics of a traditional RDBMS
- A document database allows some form of data description without enforcing a schema
- The alignment with web-development programming practices has resulted in JSON and document databases/storage

Introduction

- Let us see what are those formats and how they are used!
- We will start with eXtensible Markup Language
- Then we will look at JavaScript Object Notation and it's Binary version

eXtensible Markup Language (XML)

- Defined by the WWW Consortium (W3C)
- Extensible, unlike HTML, users can add new tags, and separately specify how the tag should be handled for display
- XML has become the basis for all new generation data interchange formats. For instance bank transfers and secure document exchange
- Documents have tags giving extra information about sections of the document. Those tags can also be nested

```
<!DOCTYPE html>
<html xmlns="http://www.w3.org/1999/xhtml">
  <head>
    <meta charset="UTF-8"/>
    <title>Polyglot (X)HTML Template</title>
  </head>
  <body>
    <p>content....</p>
  </body>
</html>
```

XML

- Each XML based standard defines what are valid elements, using XML type specification languages to specify the syntax
- DTD (Document Type Descriptors): describes the structure of an XML document
- XML Schema (newer than DTD): a special type of XML document that describes the elements that may be present
- Sample implementation database BaseX (basex.org)

Document Type Descriptors

Document
Database and
MapReduce

Introduction

NoSQL
database and
CAP theorem

NoSQL
database and
CAP theorem

NoSQL
database and
CAP theorem

RDBMS vs
Document
database

Document
Database

Document
Data-Model

MapReduce

12 / 36

```
<! ELEMENT department(dept_name, building, budget)>
<! ELEMENT dept_name (#PCDATA)>
<! ELEMENT budget (#PCDATA)>
<! ELEMENT university ( ( department | course |
    instructor | teaches )+)>
```

Notation:

	:	alternatives
+	:	1 or more occurrences
*	:	0 or more occurrences
#PCDATA	:	Parsed character data i.e. parsed string

XML Processing

- XPath: A syntax for retrieving specific elements from an XML document using wildcards.
- XQuery: A query language provides mechanisms for modifying a document. XQuery is sometimes referred to as “the SQL of XML”.
- XSLT (Extensible Stylesheet Language Transformations): A language for transforming XML documents into alternative formats, including non-XML formats such as HTML.
- DOM (Document Object Model): An object-oriented API that programs can use to interact with XML, XHTML, and similarly structured documents.

Tree Model of XML Data

- Query and transformation languages are based on a tree model of XML data
- An XML document is modeled as a tree, with nodes corresponding to elements and attributes

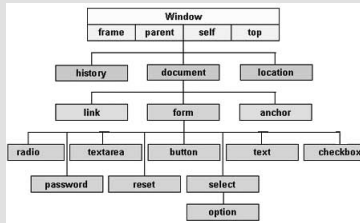


Figure: DOM sample of an HTML document

XML Processing: XPath

- XPath: is used to address (select) parts of documents using path expressions
 - The initial denotes root of the document (above the top-level tag)
 - Think of file names in a directory hierarchy
 - Selection predicates may follow any step in a path, in []
 - It is possible to apply selection criteria on the values using comparison operators ^a

^aDemonstrate an example

XML Processing: XQuery and XPath

- XQuery is derived from the Quilt query language, which itself borrows from SQL.
- XQuery uses a: for ... let ... where ... order by ... result ...^{a b}
 - for = from
 - where = where
 - order by = order by
 - result = select

^a**let**: allows temporary variables, and has no equivalent in SQL

^bDemonstrate an example

JavaScript Object Notation JSON

- JSON is an open-standard format that uses human-readable text to transmit data objects consisting of attribute–value pairs
- JSON Schema is based on the concepts from XML Schema, but is JSON-based
- Document databases use JSON documents in order to store records, just as tables and rows store records in a relational database

BSON

- BSON: binary-encoded format used in MongoDB instead of JSON
- BSON extends the JSON model to provide additional data types such as integer and float to be efficient for encoding and decoding within different languages.
- BSON implementation supports embedding objects and arrays within other objects and arrays

MongoDB

- A MongoDB instance may have zero or more databases
- A database may have zero or more collections.
 - Can be thought of as the relation (table) in DBMS, but with many differences.
- A collection may have zero or more documents.

MongoDB

- A MongoDB instance may have zero or more databases
- A database may have zero or more collections.
 - Can be thought of as the relation (table) in DBMS, but with many differences.
- A collection may have zero or more documents.
 - Docs in the same collection don't even need to have the same fields
 - Docs are the records in RDBMS
 - Docs can embed other documents
 - Documents are addressed in the database via a unique key differences.

MongoDB

- A MongoDB instance may have zero or more databases
- A database may have zero or more collections.
 - Can be thought of as the relation (table) in DBMS, but with many differences.
- A collection may have zero or more documents.
 - Docs in the same collection don't even need to have the same fields
 - Docs are the records in RDBMS
 - Docs can embed other documents
 - Documents are addressed in the database via a unique key differences.
- A document may have one or more fields.
- query language is JavaScript
- There is no join provided in MongoDB. You have to implement it manually.

Document
Database and
MapReduce

Introduction

NoSQL
database and
CAP theorem

NoSQL
database and
CAP theorem

NoSQL
database and
CAP theorem

RDBMS vs
Document
database


Document
Database

Document
Data-Model

MapReduce

Relational Model

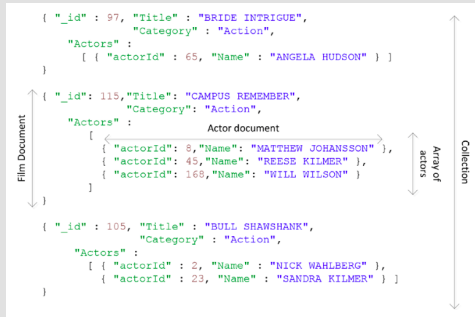
- Suppose you have the following entities and their relationships



- How would we model this in a document structure?

Embedded

- First mapping possibility is to map to one embedded collection.



- document database are not designed to be normalized and data repetition is accepted, but could have side effects.

Linking using object _id

- Second mapping possibility is to map to different collections and link the documents

```
{ "_id": 115, "Title": "CAMPUS REMEMBER", "Category": "Action",  
  "Actors": [8, 45, 168]}
```

```
{ "actorId": 168, "Name": "WILL WILSON" }
```

```
{ "actorId": 45, "Name": "REESE KILMER" }
```

```
{ "actorId": 8, "Name": "MATTHEW JOHANSSON" }
```

- This approach is less suited for document databases since the binary data of those collections are not stored as a continuous stream.
- Another disadvantage of this approach is the lack of join query

Querying collections and objects

- Select queries in MongoDB

\\SQL

```
SELECT * FROM actors
```

```
SELECT * FROM actors WHERE age = 23
```

```
SELECT * FROM actors WHERE age = 23 ORDER BY name
```

\\Mongo

```
db.actors.find()
```

```
db.actors.find({age: 23})
```

```
db.actors.find({age: 23}).sort({name:1})
```

Querying collections and objects

- Insert queries in MongoDB
- Suppose we have a relationship between actor entity and address entity based on actor_id

\\SQL

```
INSERT INTO actors(actor_id,name,age)
          VALUES(3, "actor name", 45)
INSERT INTO address(addressid, street, city,
                    actor_id)
          VALUES(5, "Wijnhaven 66", "Rotterdam", 3)
```

\\Mongo

```
db.actors.insert({name:"actor name", age: 23,
                  address:{street:"Wijnhaven 66",
                           city:"Rotterdam"}}
})
```

Querying collections and objects

- Update queries in MongoDB

\\SQL

```
UPDATE actors
  SET name = "New name"
  WHERE actor_id = "a1"
```

\\Mongo

```
db.actors.update(
  { _id: "a1" },
  \\& sould be dollar
  { &set : {
                                name: "New Name"}
  })
```

Introduction


- MapReduce is a data processing paradigm for condensing large volumes of data into useful aggregated results.
- Map- and Reduce functions are commonly used in functional programming
- In INFDEV02-2 and INDEV02-3 we already introduced HOFs
- MapReduce rely on the concept of higher order functions HOFs are very powerful in the context of NoSQL databases.
- The following functions will be further discussed : FlatMap, Map and Reduce

Map Function

- Apply the function f to each element of list x
- `map(f, x[0...n-1])`
- in Python:

```
def square(x):  
    return x * x
```


```
map(square, [1, 2, 3, 4]) #would return [1, 4, 9,  
16]
```



Reduce Function

- Repeatedly apply binary function f to pairs of items in x , replacing the pair of items with the result until only one item remains
- `reduce(f, x[0...n-1])`
- in Python:

```
def add(x, y):
    return x+y
reduce(add, [1,2,3,4]) #would result in a 10
```



FlatMap Function

Document
Database and
MapReduce

Introduction

NoSQL
database and
CAP theorem

NoSQL
database and
CAP theorem

NoSQL
database and
CAP theorem

RDBMS vs
Document
database

Document
Database

Document
Data-Model

MapReduce

29 / 36

- Repeatedly apply function f to items in a sublist, then removing the sublist structure with the result of one dimensional list
- $\text{FlatMap}(f, x[[0..n-1],[0..m-1]])$
- in pseudo Python: suppose the f function returns the element without any change

```
listOfLists = [[1, 2],[3, 4, 5], [6]]  
for l in listOfLists:  
    map(f, l)  
reduce(list.__add__, listOfLists)  
#would result in a flatten list [1, 2, 3, 4, 5, 6]
```

SQL and MapReduce


- We have seen so far what map en reduce functions are
- But why do we need them in document sturcture?
- What are the similarities between MapReduce and SQL?

MapReduce Function in MongoDB

- Data in mongoDB are saved in documents
- The MapReduce function first queries the collection, then maps the result documents to emit key-value pairs which is then reduced based on the keys that have multiple values.

```

Collection
  ↓
db.orders.mapReduce(
  map   → function() { emit( this.cust_id, this.amount ); },
  reduce → function(key, values) { return Array.sum( values ); },
  query → { status: "A" },
  output → "order_totals"
)
  
```



Map and Filter functions vs SQL

- Relational databases use the map, filter and reduce paradigm (where it is called project, select, aggregate).
- `SELECT MAX(pixels) FROM cameras WHERE brand = 'Nikon'`
 - cameras is a sequence (a list of rows, where each row has the data for one camera)
 - `WHERE brand = 'Nikon'` is a **filter**
 - pixels is a **map** (extracting just the pixels field from the row)
 - MAX is a **reduce**
- Demo!

Join as MapReduce in MongoDB

- MongoDB does not provide explicit join queries to join two collections
- To implement a join you need
 - a mapper for each collection to retrieve key and values for each collection
 - A reducer function to reduce values for each key
- Demo!

MapReduce Function in a Cluster

- How does CAP theorem effect the implementation of MapReduce?
- Generally speaking It depends on the execution of MapReduce whether local or in a cluster
- We have seen how MapReduce is executed locally in document database
- What about clusters?

MapReduce Function in a Cluster

- The distributed MapReduce idea is similar to (but not the same as!): $\text{reduce}(f2, \text{map}(f1, x))$
- Key idea: "data-centric" architecture – Send function $f1$ directly to the data : Execute it concurrently
- Then merge results with reduce: Also concurrently

Document
Database and
MapReduce

Introduction

NoSQL
database and
CAP theorem

NoSQL
database and
CAP theorem

NoSQL
database and
CAP theorem

RDBMS vs
Document
database

Document
Database

Document
Data-Model

MapReduce

End

- Thank you and the best of luck!